

# Machine & Deep Learning

Pr Rahhal ERRATTAHI

[rahhal.errattahi@um6p.ma](mailto:rahhal.errattahi@um6p.ma)

[errattahi.r@ucd.ac.ma](mailto:errattahi.r@ucd.ac.ma)

Lecture 01

# Course Topics

- Part 1: Introduction to ML
- Part 2: Supervised ML problem setup and Data Preprocessing
- Part 3: Kernel Methods
- Part 4: Decision Trees
- Part 5: Regression
- Part 6: Ensemble learning
- Part 7: Supervised learning of Neural Networks

# Lab sessions

- Python notebooks in CoLab
- Alternatively, set up Jupyter Notebooks yourself
  - Numpy
  - Matplotlib
  - scikit learn
  - tensorflow
  - PyTorch

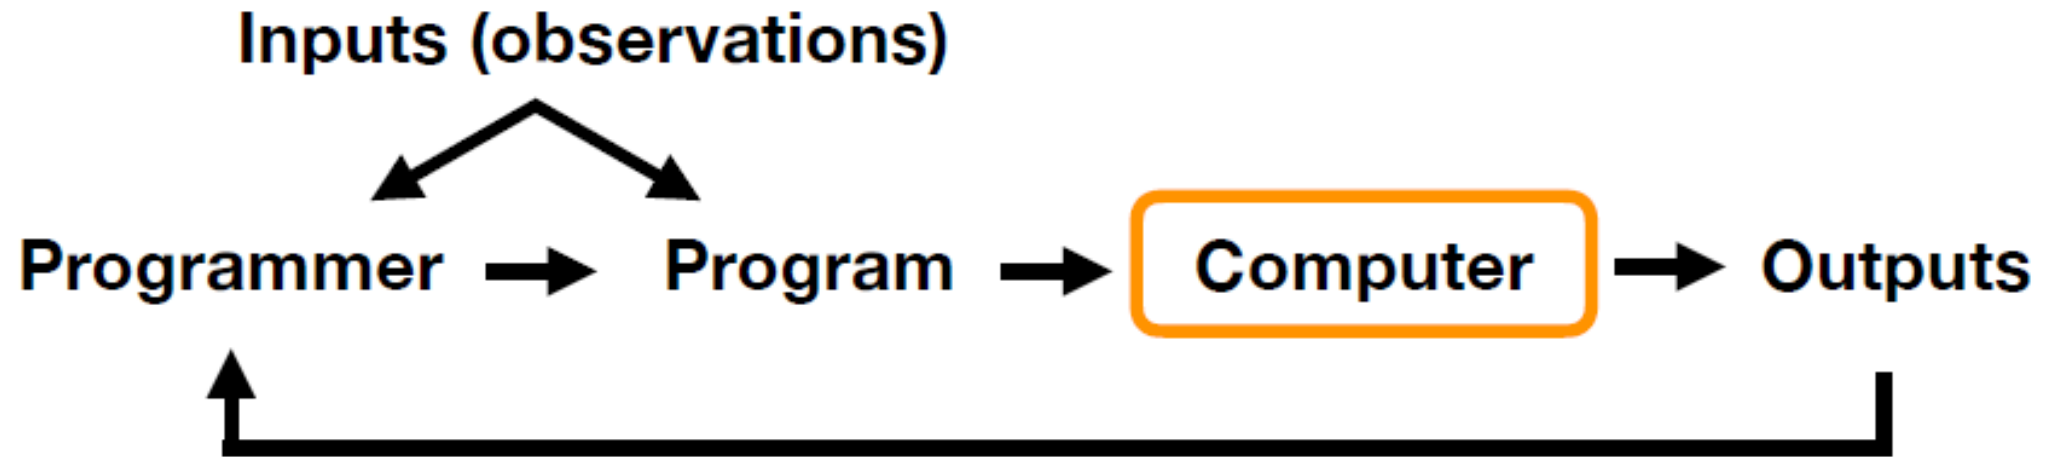
# Lecture 1 Overview

- What is machine learning ?
- Categories of machine learning
- Notation
- Approaching a machine learning application
- ML Terminology

# Lecture 1 Overview

- What is machine learning ?
- Categories of machine learning
- Notation
- Approaching a machine learning application
- ML Terminology

# The Traditional Programming Paradigm



# What is Machine Learning?

“Machine learning is the field of study that gives computers the ability to learn without being explicitly programmed”

-- Arthur L. Samuel, AI pioneer, 1959



# What is Machine Learning?

- Machine learning algorithms are algorithms that **learn** models from data / experience.
- No need to formulate explicit rules.
- Algorithm performance gets better with experience / data.






# Defining the Learning Task

“A computer program is said to learn from experience **E** with respect to some class of tasks **T** and performance measure **P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**.”

— Tom Mitchell, Professor at Carnegie Mellon University

## Handwriting Recognition Example:



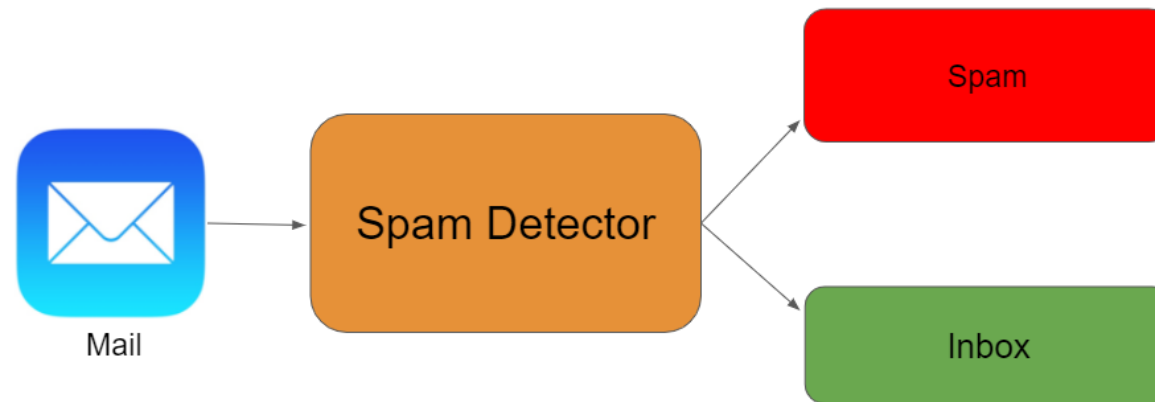
Task **T**: ?

Performance measure **P**: ?

Training experience **E**: ?

# Defining the Learning Task

Improve on task T, with respect to performance metric P, based on experience E




T: Categorize email messages as spam or legitimate.

P: Percentage of email messages correctly classified.

E: Database of emails, some with human-given labels

# Defining the Learning Task

Improve on task T, with respect to performance metric P, based on experience E



T: Playing checkers

P: Percentage of games won against an arbitrary opponent

E: Playing practice games against itself

# Defining the Learning Task

Improve on task T, with respect to performance metric P, based on experience E



T: Driving on four-lane highways using vision sensors

P: Average distance traveled before a human-judged error

E: A sequence of images and steering commands recorded while observing a human driver.

# When Do We Use Machine Learning?

ML is used when:

- Human expertise does not exist (navigating on Mars)
- Humans can't explain their expertise (speech recognition)
- Models must be customized (personalized medicine)
- Models are based on huge amounts of data (genomics)

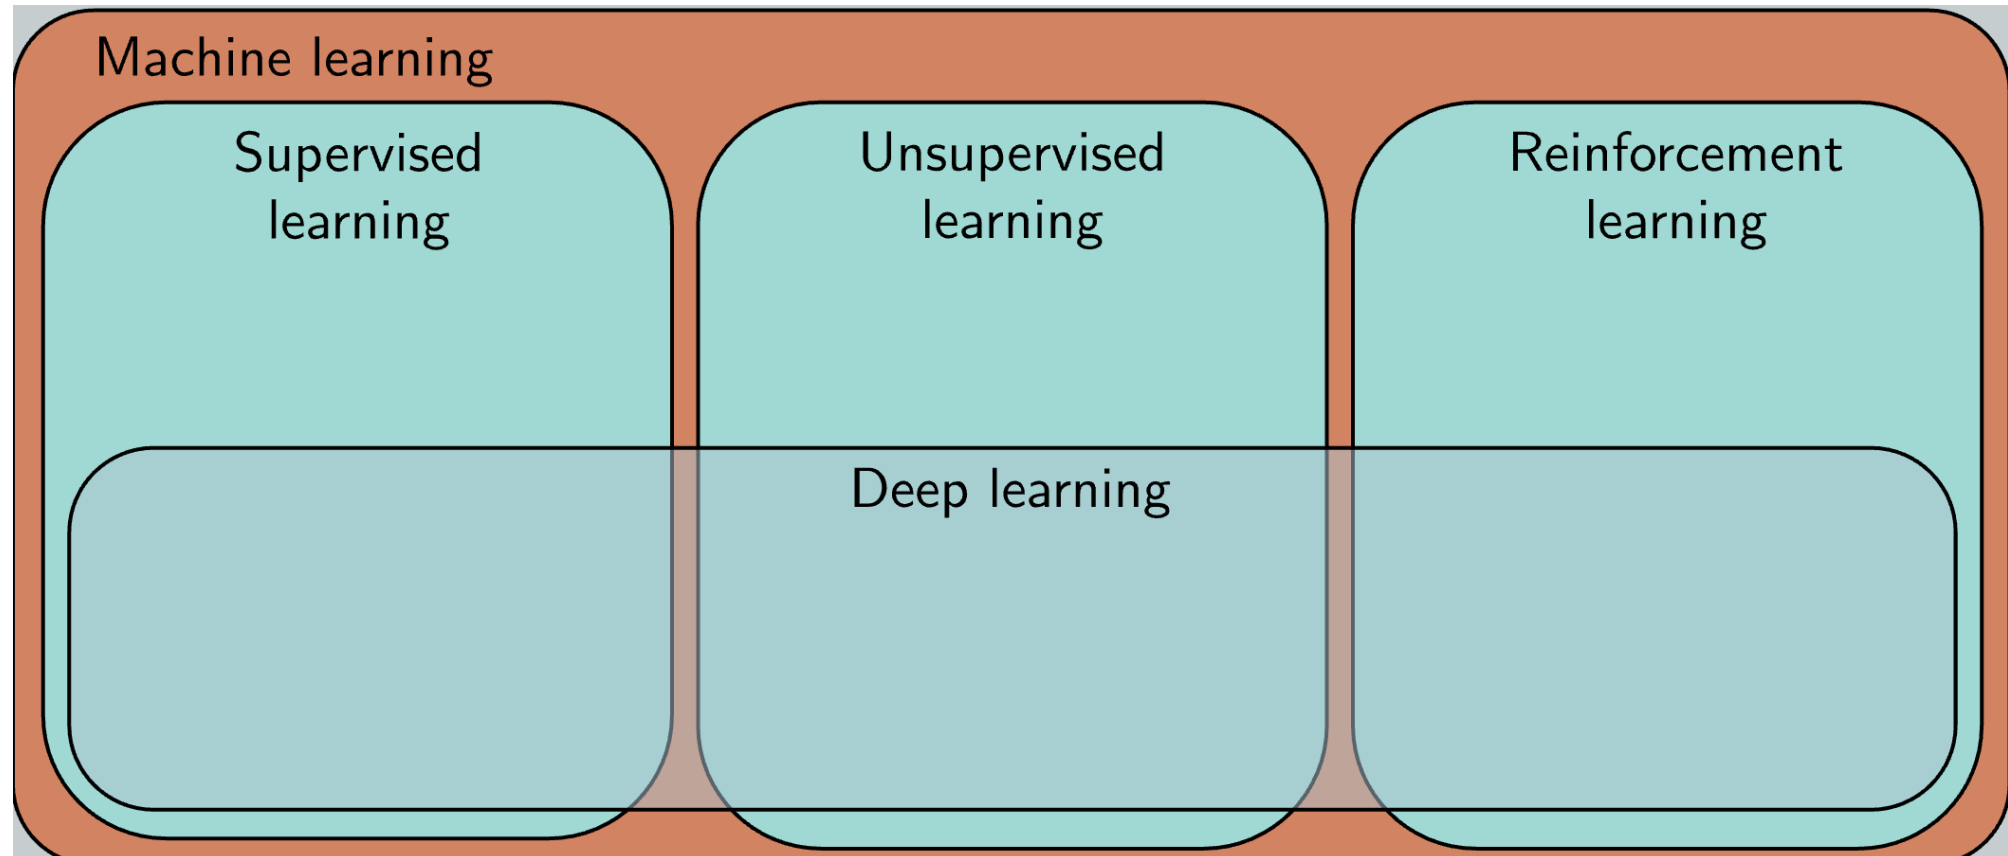
Learning isn't always useful:

- There is no need to “learn” to calculate payroll

# Lecture 1 Overview

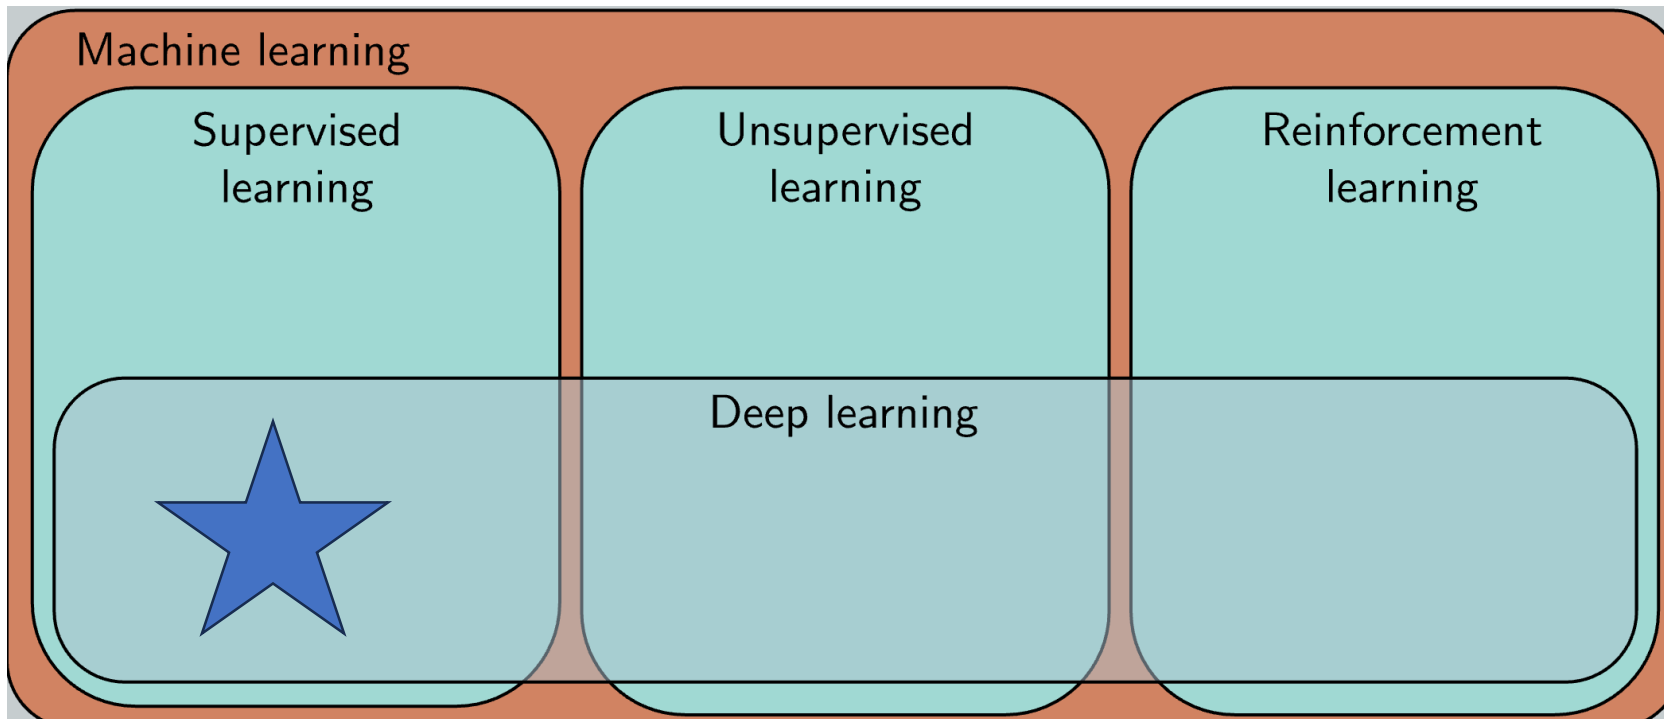
- What is machine learning ?
- Categories of machine learning
- Notation
- Approaching a machine learning application
- ML Terminology

# Categories of machine learning



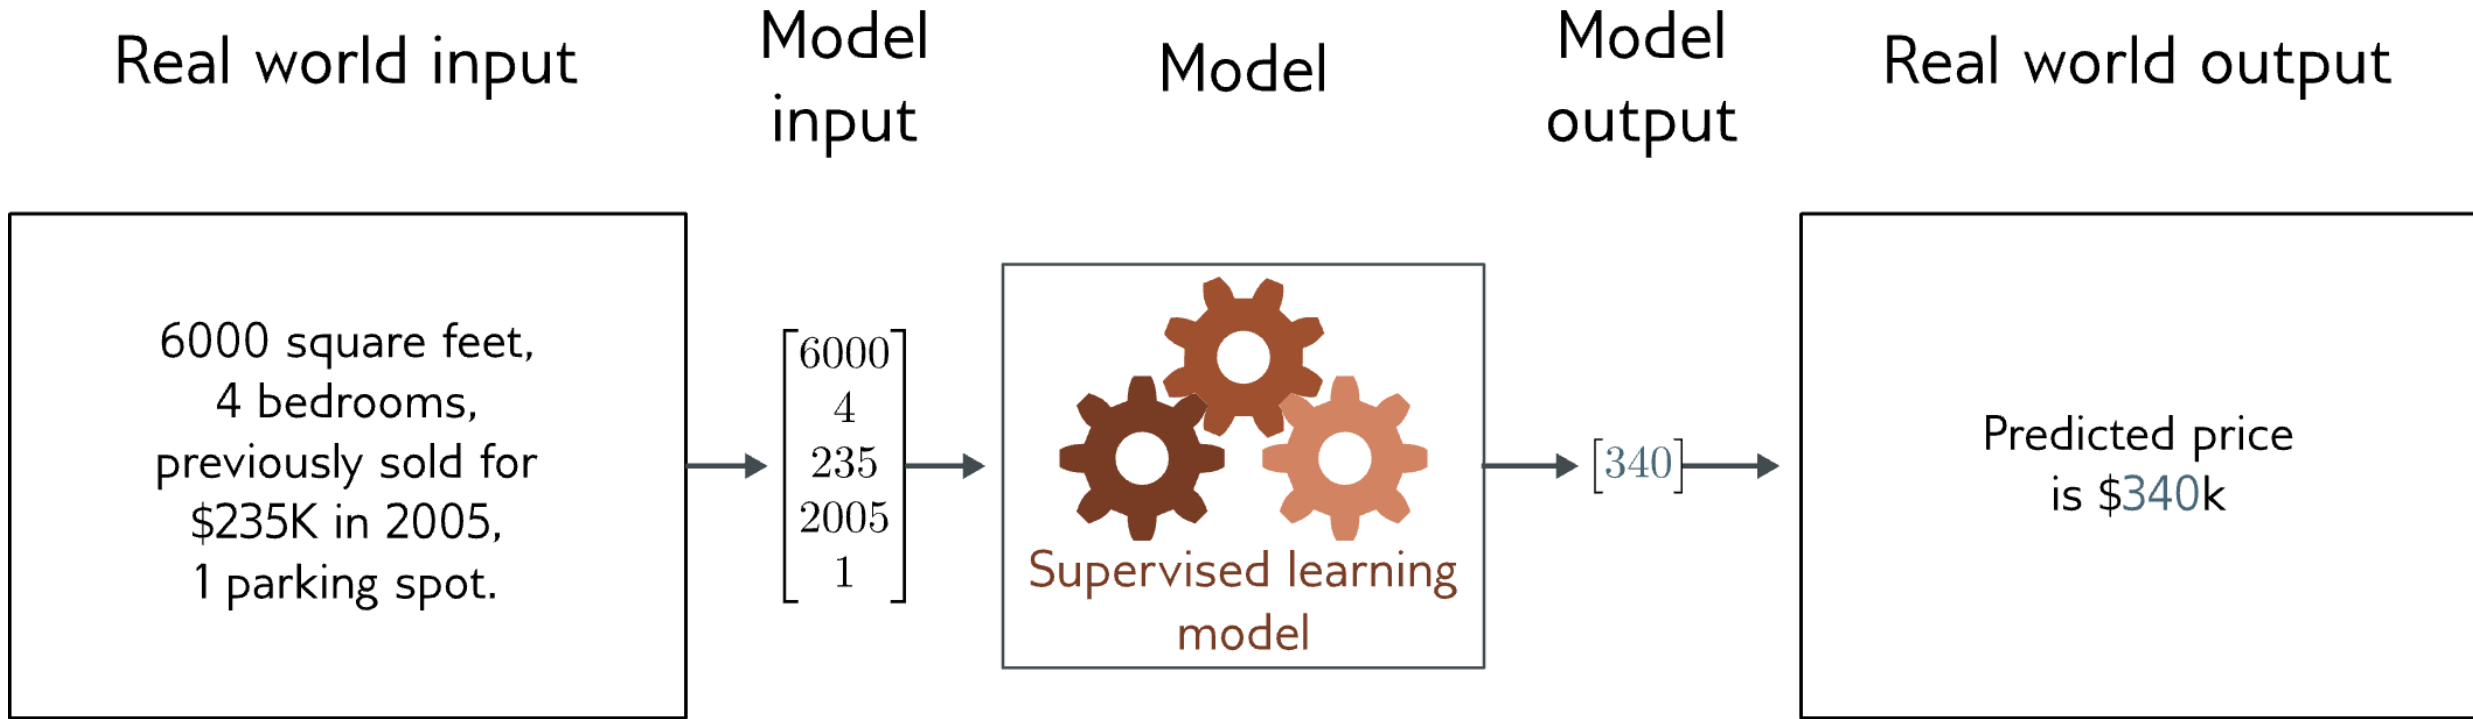
# Supervised learning

- Define a mapping from input to output
- Learn this mapping from paired input/output data examples



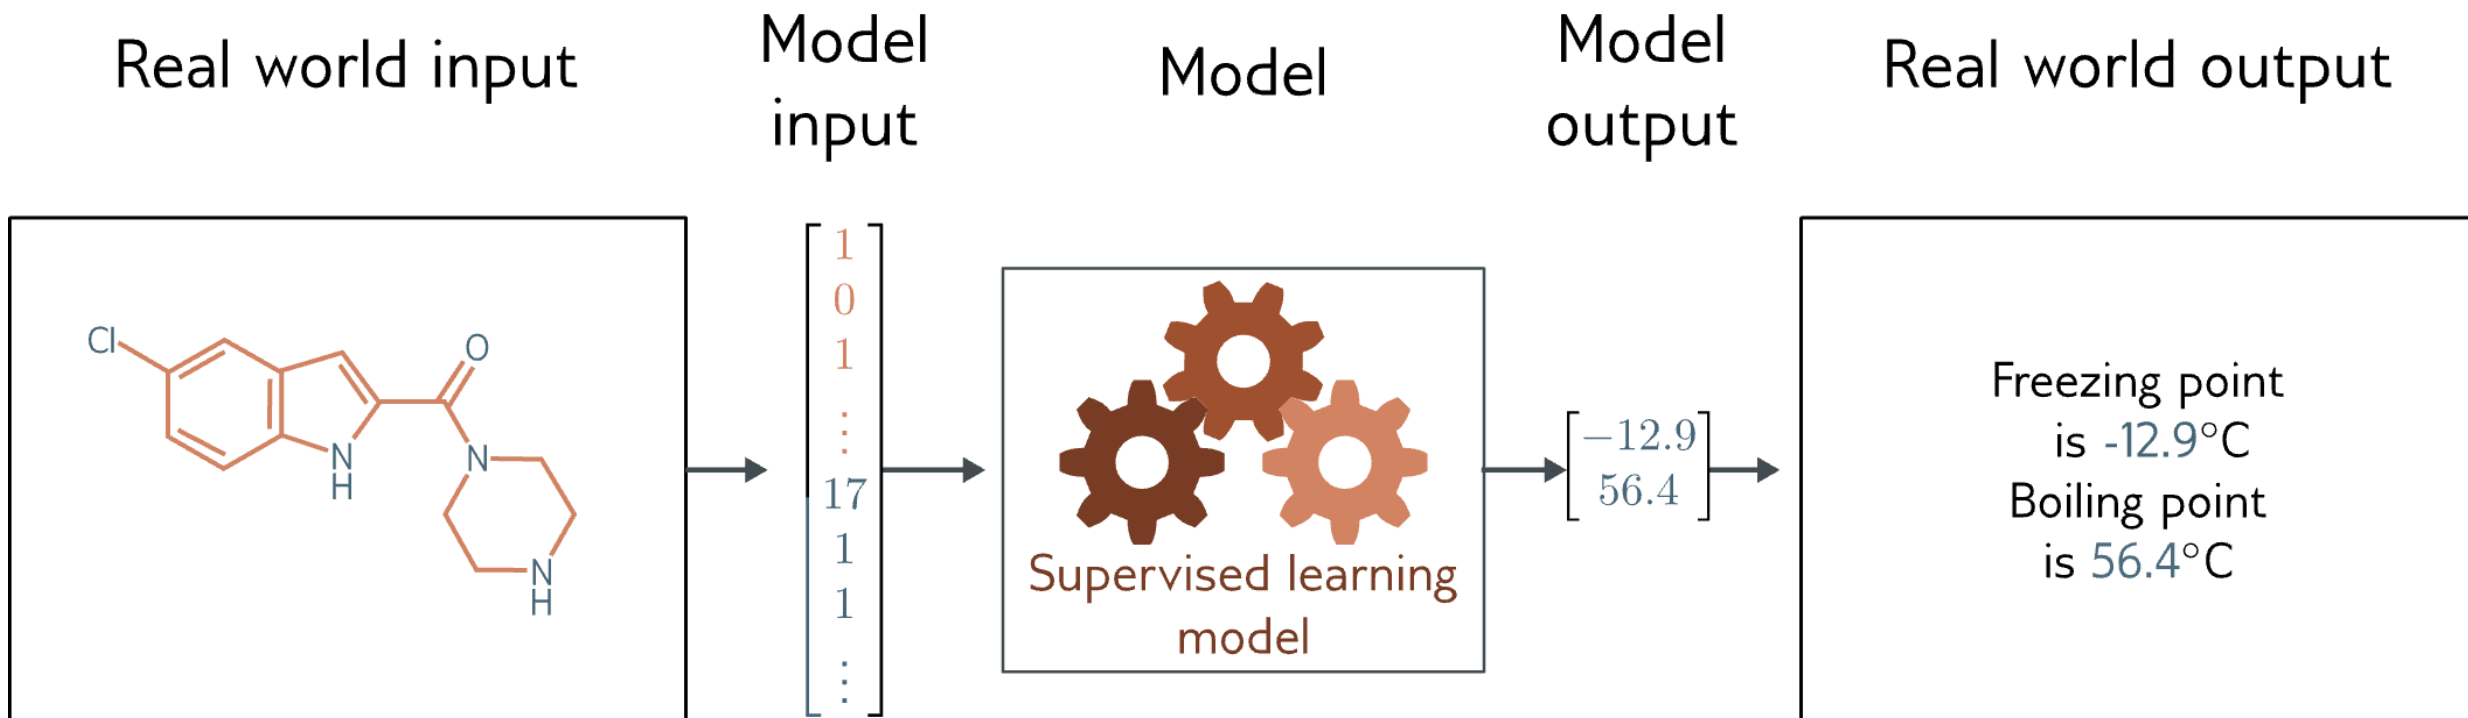


# Regression



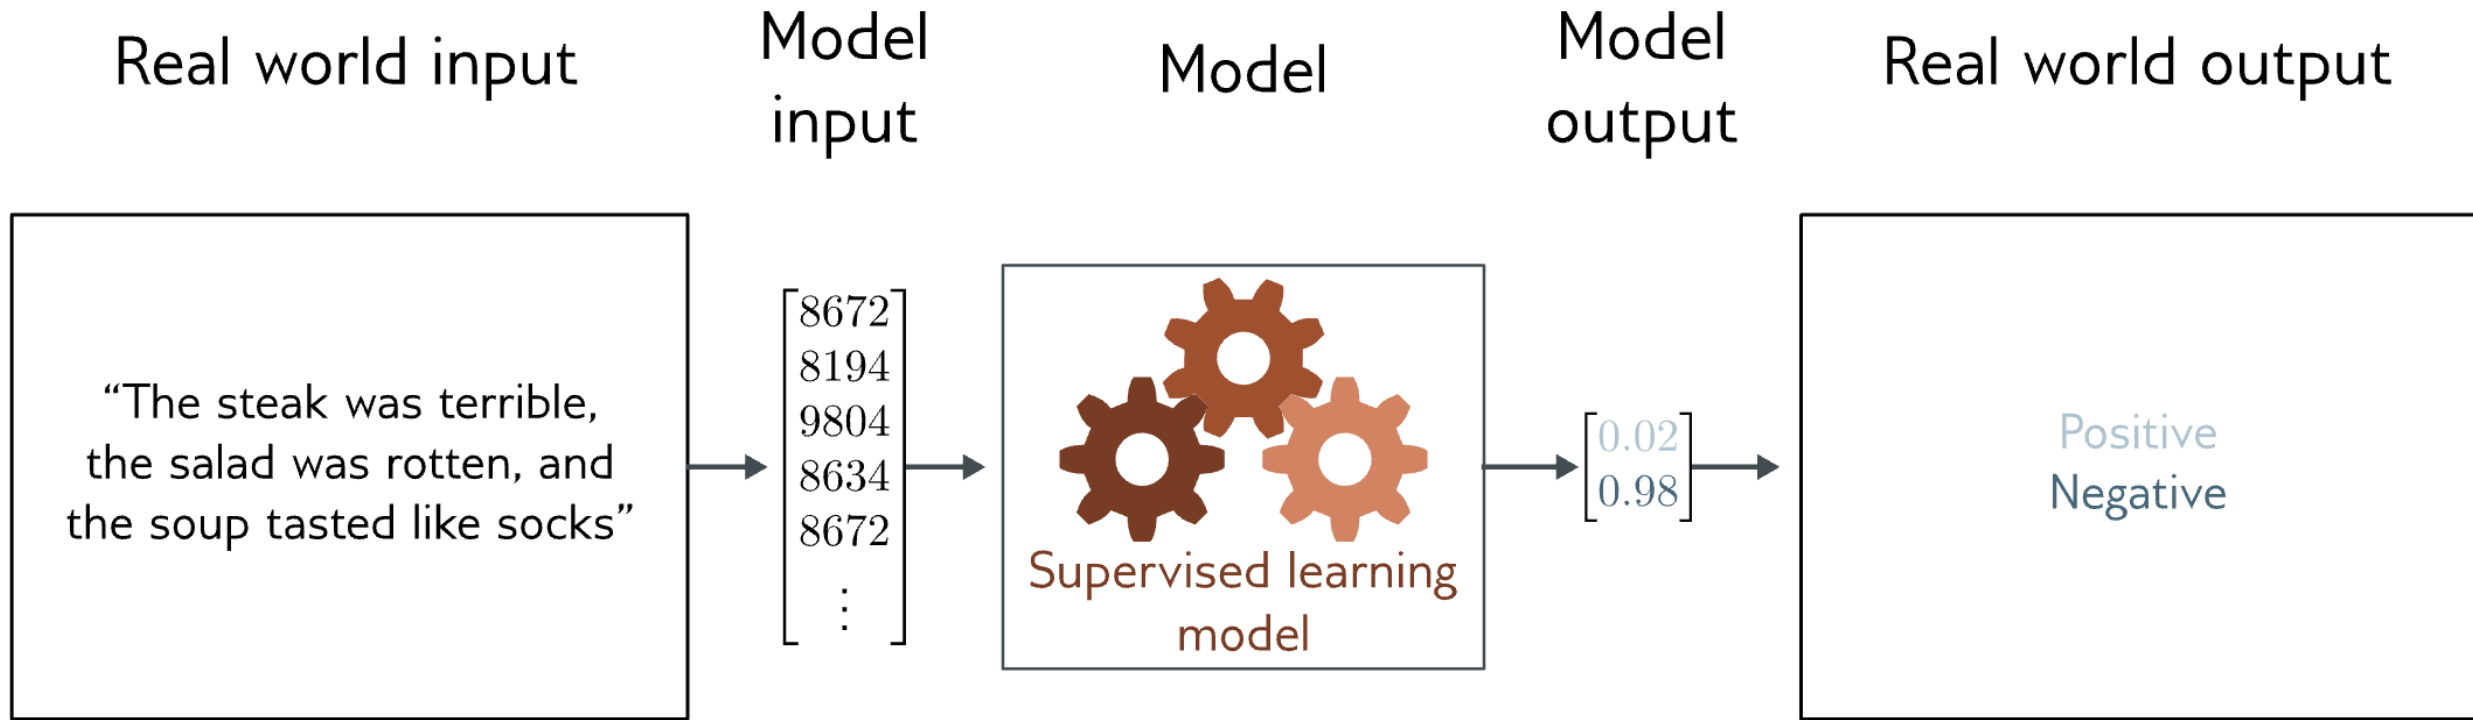
- Univariate regression problem (one output, real value)
- Fully connected network

# Graph regression



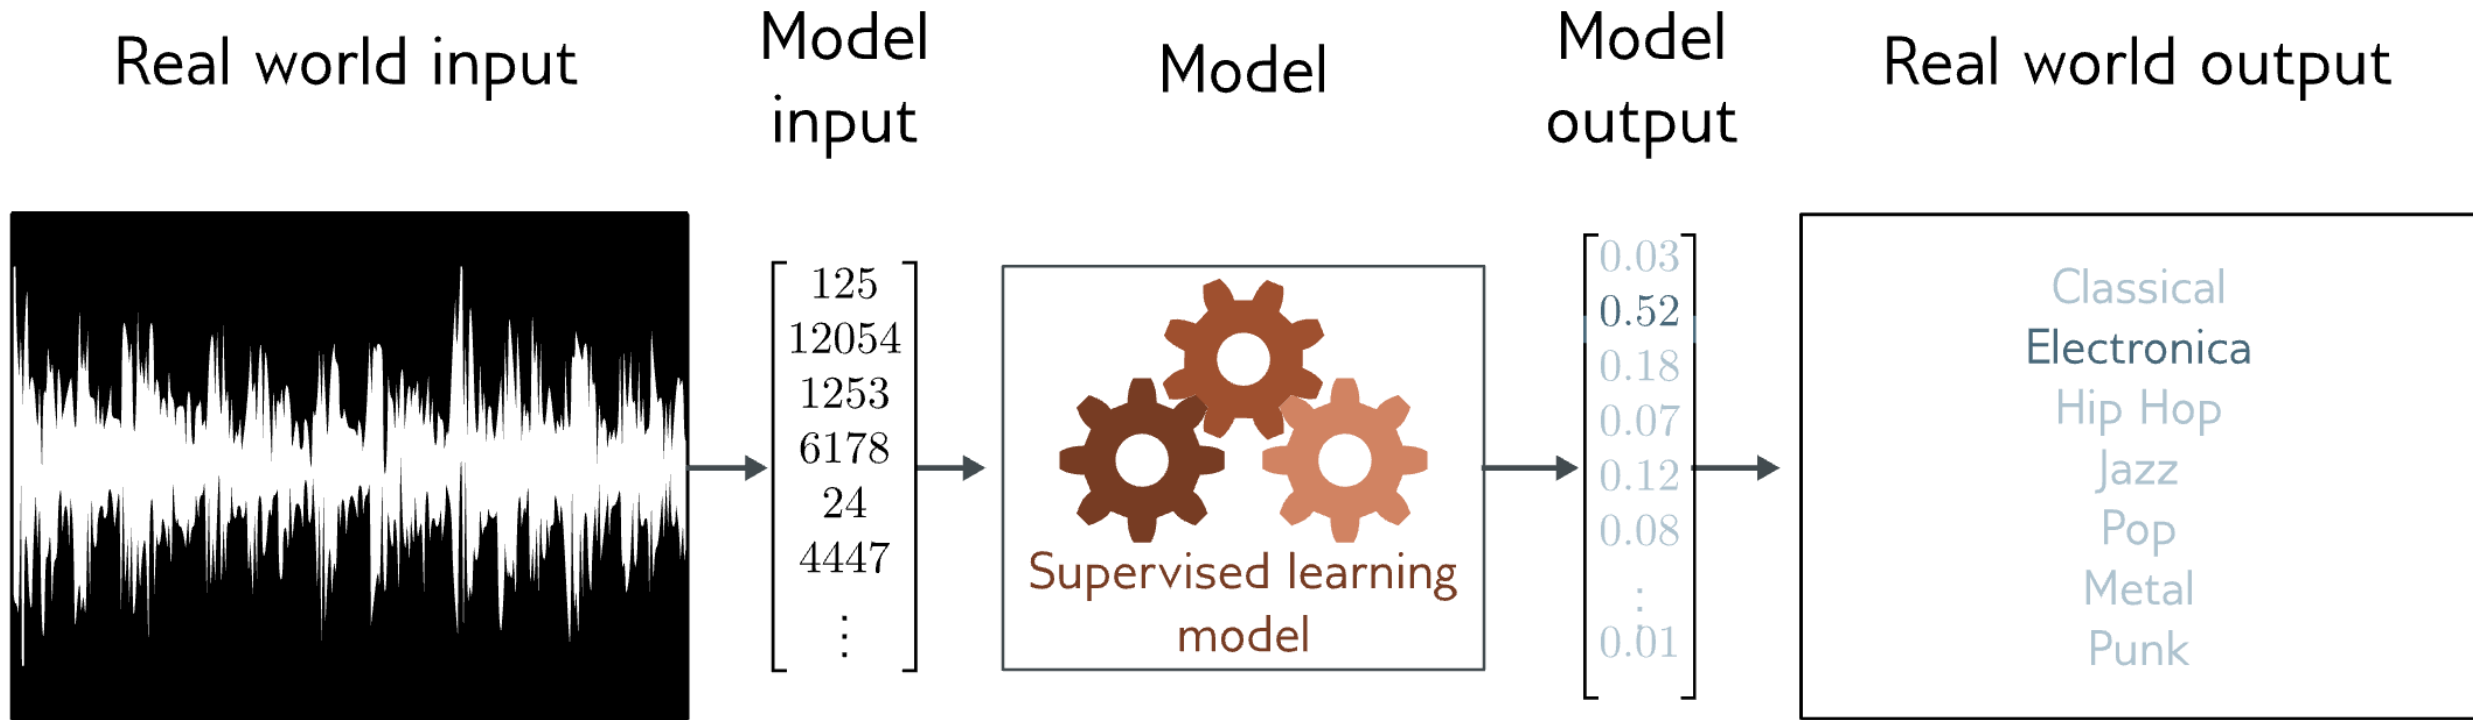
- Multivariate regression problem (>1 output, real value)
- Graph neural network

# Text classification



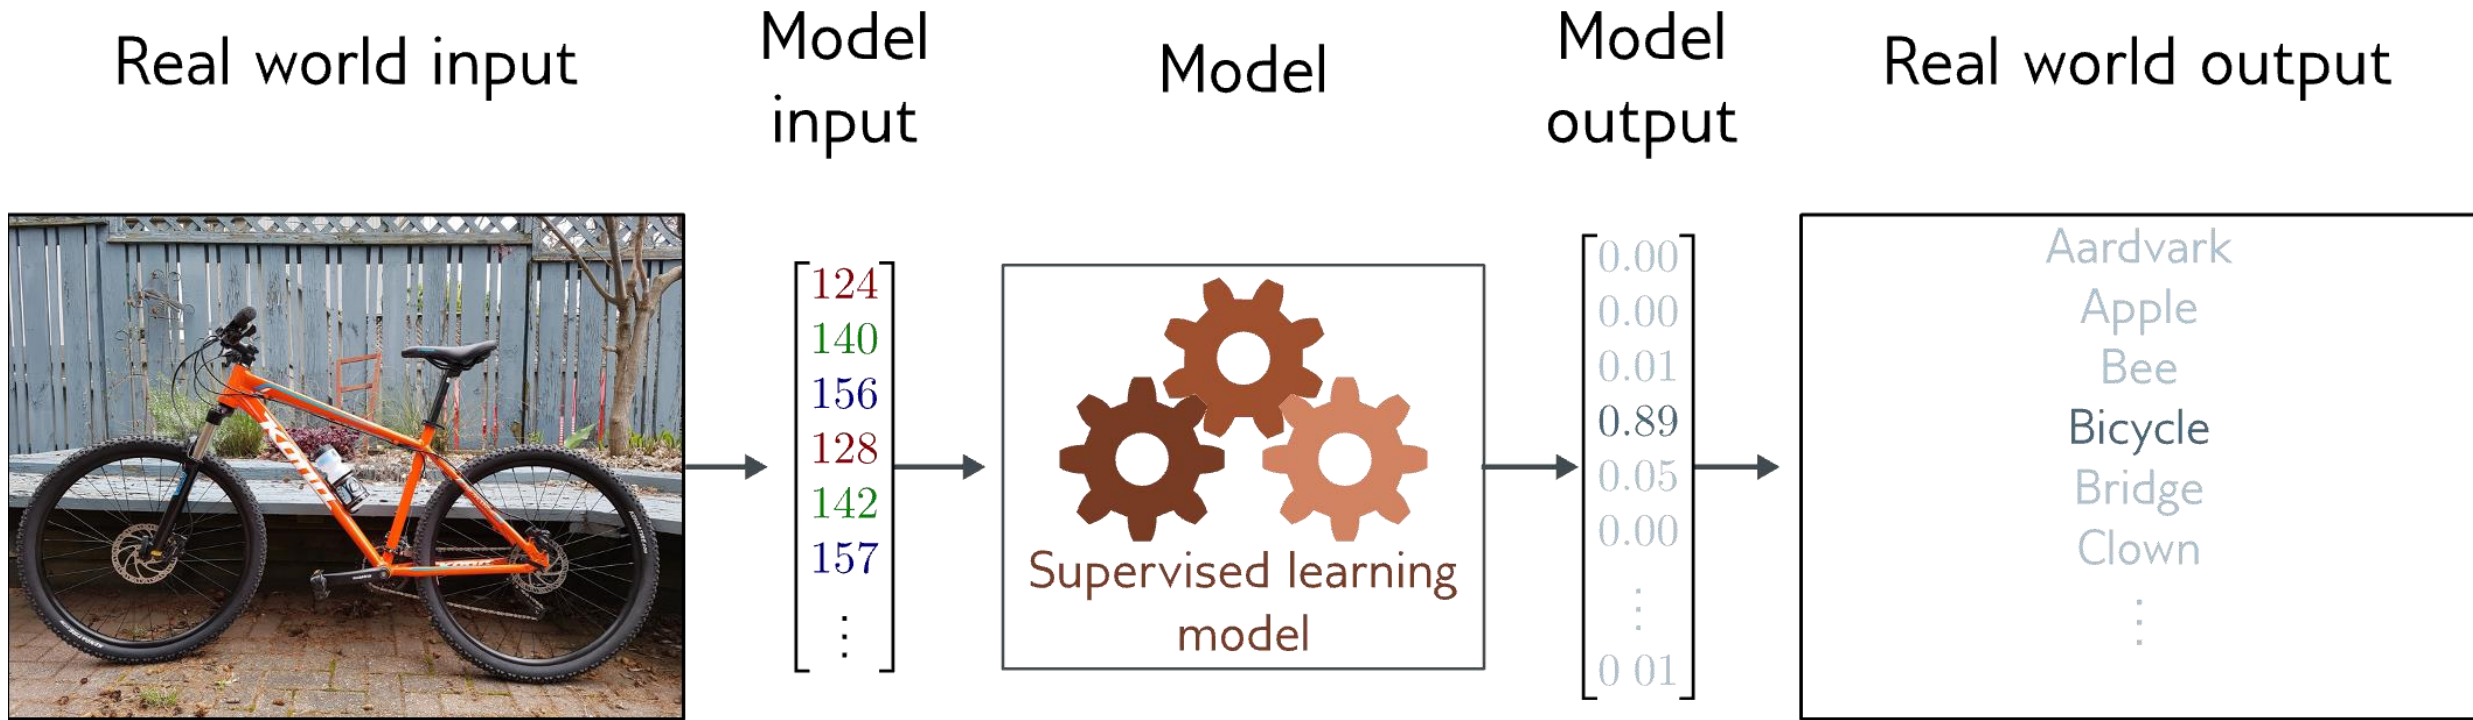
- Binary classification problem (two discrete classes)
- Transformer network

# Music genre classification




- Multiclass classification problem (discrete classes, >2 possible values)
- Recurrent neural network (RNN)

# Image classification



- Multiclass classification problem (discrete classes, >2 possible classes)
- Convolutional network

# Image segmentation

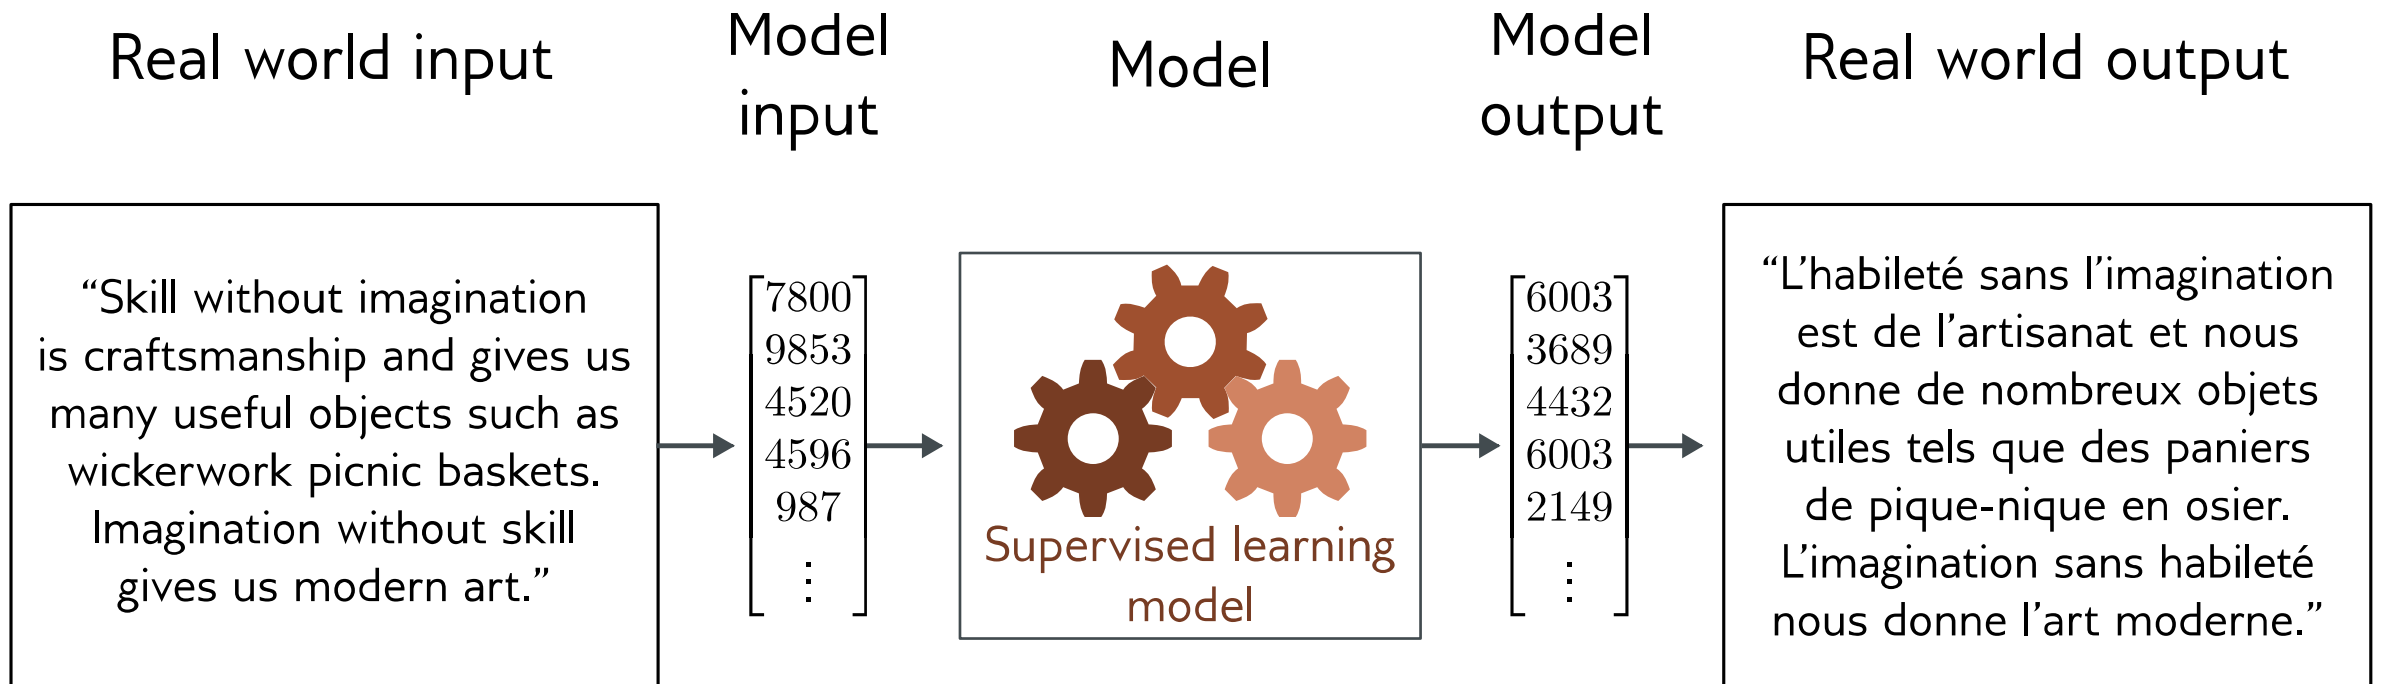


- Multivariate binary classification problem (many outputs, two discrete classes)
- Convolutional encoder-decoder network

# ML Terminology

- **Regression** = continuous numbers as output
- **Classification** = discrete classes as output
- **Two class** and **multiclass** classification treated differently
- **Univariate** = one output
- **Multivariate** = more than one output

# Translation





# Image captioning


Real world input



Model  
input

$$\begin{bmatrix} 183 \\ 204 \\ 231 \\ 185 \\ 204 \\ 232 \\ \vdots \end{bmatrix}$$

Model



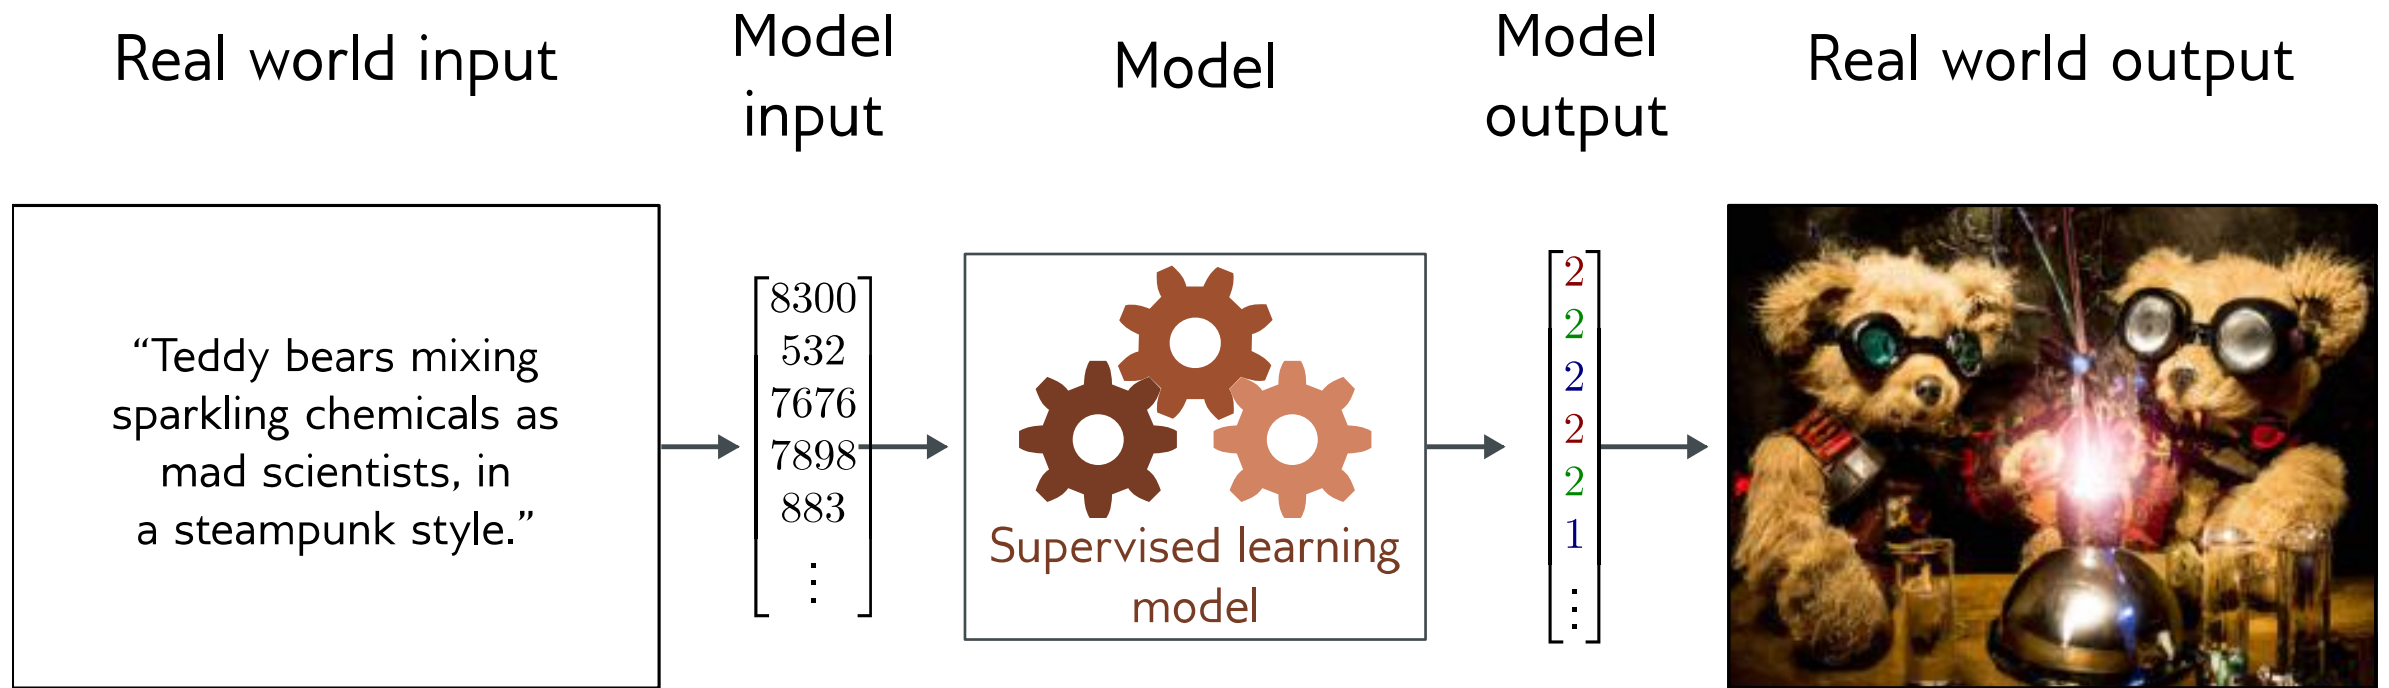
Model  
output

$$\begin{bmatrix} 1 \\ 5593 \\ 7532 \\ 7924 \\ 1 \\ \vdots \end{bmatrix}$$

Real world output

"A Kazakh man on a  
horse holding a  
bird of prey"

# Image generation from text



# What do these examples have in common?

- Very complex relationship between input and output
- Sometimes may be many possible valid answers
- But outputs (and sometimes inputs) obey rules

“A Kazakh man on a  
horse holding a  
bird of prey”

Language obeys  
grammatical rules



Natural images also  
have “rules”

# Supervised learning

Supervised learning is about function approximation

## Problem Setting:

- Set of possible instances  $X$
- Unknown target function  $f : X \rightarrow Y$
- Set of function hypotheses  $H = \{h \mid h : X \rightarrow Y\}$

## Input:


- training examples  $\{ \langle x_i, y_i \rangle \}$ . For example  $x$  is an email and  $y$  is either Spam or No Spam.

## Output:

- Hypothesis  $h \in H$  that best approximates target function  $f$ . OR
- a classification “rule” that can determine the class of any object from its attributes values.


# Unsupervised Learning

- Unsupervised learning is about description, opposed to approximation (supervised learning).
  - Clustering
  - Finding outliers
  - Generating new examples
  - Filling in missing data










DeepCluster: Deep Clustering for Unsupervised Learning of Visual Features (Caron et al., 2018)



# Unsupervised Learning


- Learning about a dataset without labels
  - e.g., clustering
- Generative models can create examples
  - e.g., generative adversarial networks






# Unsupervised Learning

- Learning about a dataset without labels
  - e.g., clustering
- Generative models can create examples
  - e.g., generative adversarial networks
- PGMs learn distribution over data
  - e.g., variational autoencoders,
  - e.g., normalizing flows,
  - e.g., diffusion models



# Generative models



 National Geographic  
Domestic cat



 Wikipedia  
Cat - Wikipedia



 The Guardian  
pet guru Yuki Hattori explain | ...



 Britannica  
Cat | Breeds & Facts | Britannica

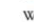


 The Spruce Pets  
Tabby Cat: Breed Profile ...



 Britannica  
Cat | Breeds & Facts | Britannica



 Wikipedia  
Cat intelligence - Wikipedia




 Smithsonian Magazine  
Cats React to 'Baby Talk' From Their ...




 Alley Cat Allies  
The Natural History of Domestic Cats ...




 The New York Times  
How the Cat Gets Its Stripe...




 Country Living Magazine  
Friendliest Cat Breeds Tha...



 Freepik  
Cat Images - Free D...



 BBC Science Focus  
What's the longest a cat can live for ...



 National Geographic  
Domestic cat



 DK Find Out!  
Cat Facts for Kids | What is a Cat | DK ...



 The Spruce Pets  
Ragdoll Cat: Breed Profile ...



 Good Housekeeping  
25 Best Cat Instagram Caption...



 Daily Paws  
17 Long-Haired Cat Breeds to Swoon...



 Unsplash  
500+ Domestic Cat ...

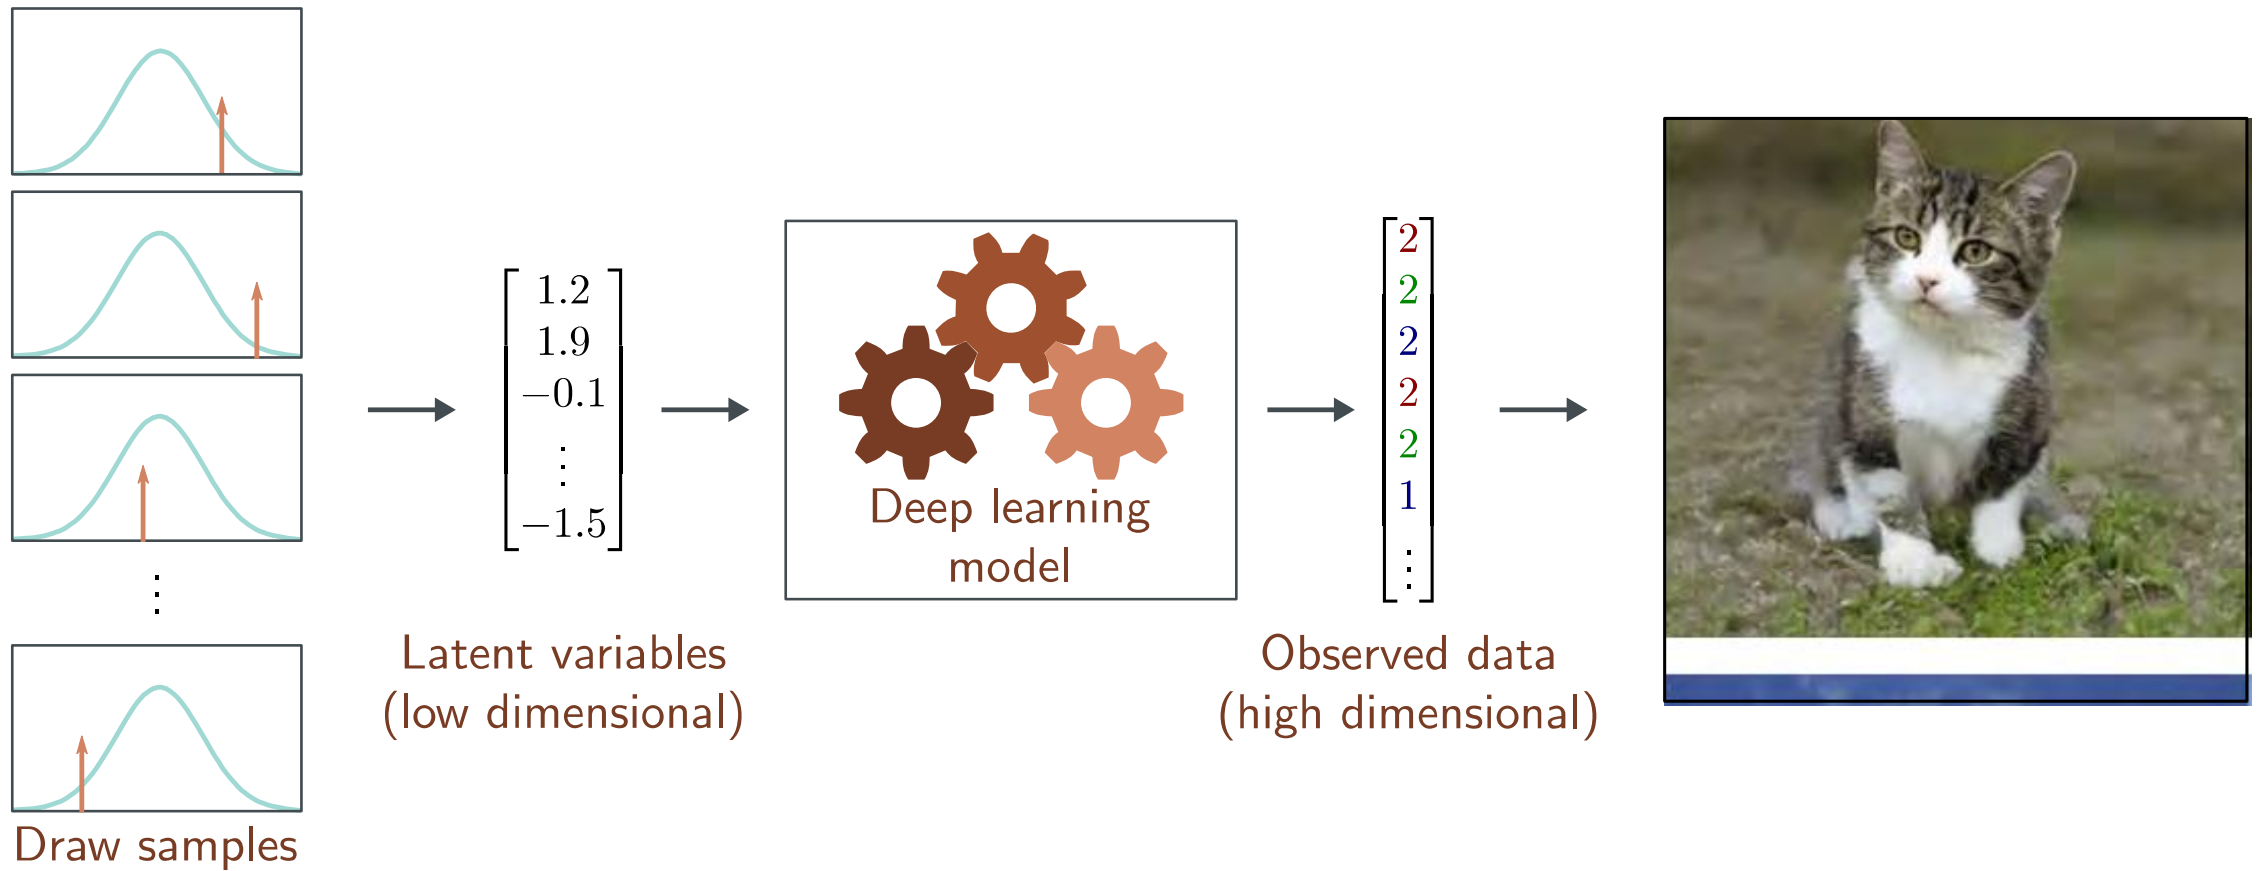


 Four Paws  
A Cat's Personality - FOUR PAWS ...



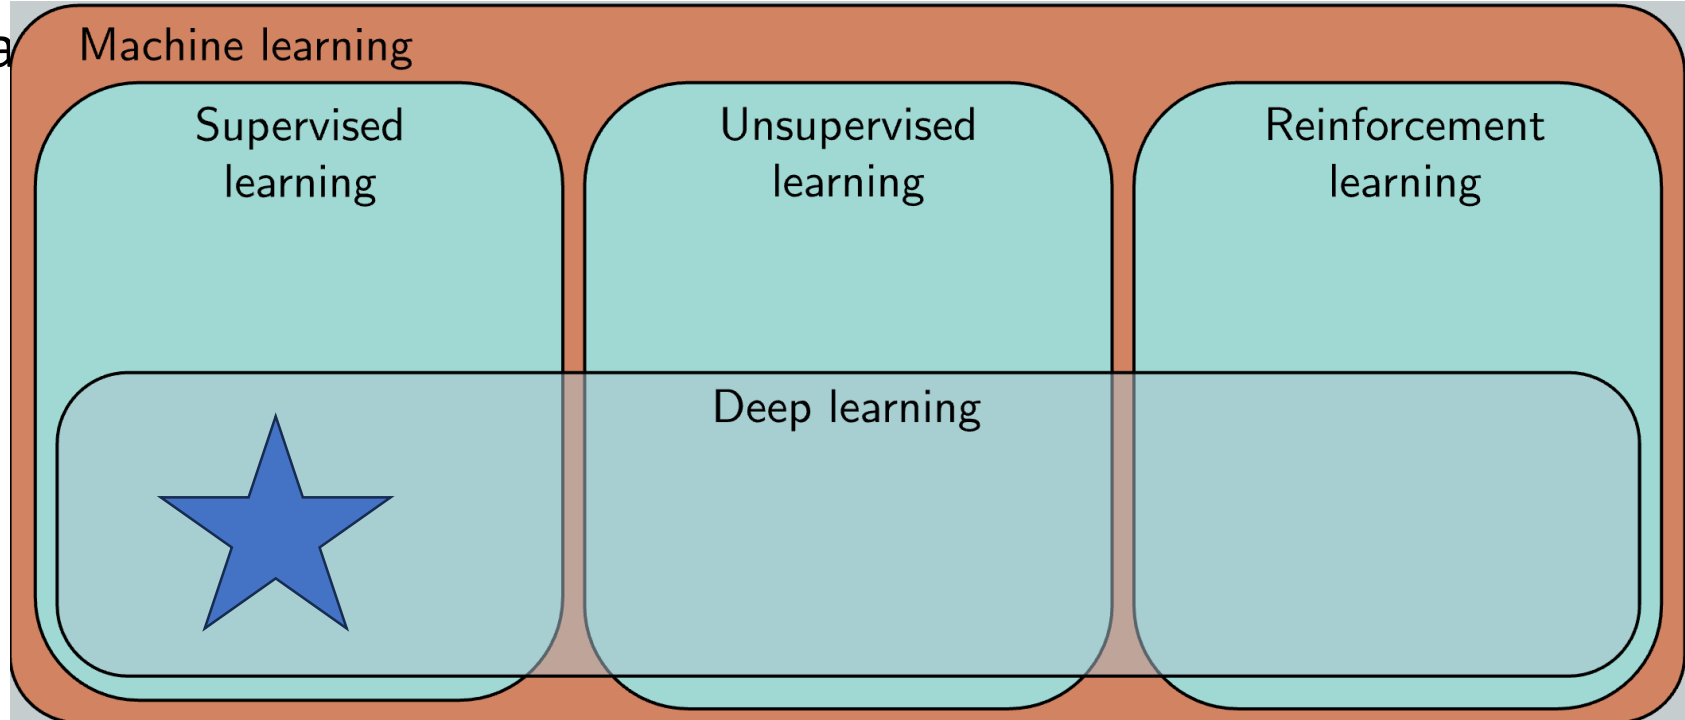
 The Guardian  
pet guru Yuki Hattori explain | ...

# Latent variables



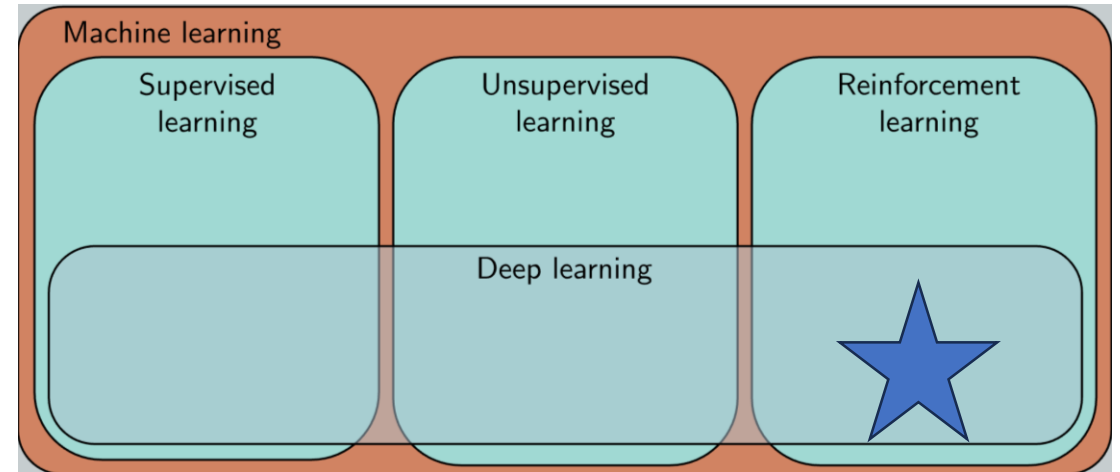
# Unsupervised Learning

- Learning about a dataset without labels
  - Clustering
  - Finding outliers
  - Generating new examples
  - Filling in missing data



# Reinforcement learning

- A set of **states**
- A set of **actions**
- A set of **rewards**



- Goal: take actions to change the state so that you receive rewards
- You don't receive any data – you have to explore the environment yourself to gather data as you go


# Example: chess

- States are valid states of the chess board
- Actions at a given time are valid possible moves
- Positive rewards for taking pieces, negative rewards for losing them

State




: Action



# Example: chess

- States are valid states of the chess board
- Actions at a given time are valid possible moves
- Positive rewards for taking pieces, negative rewards for losing them



# Why is this difficult?


- Stochastic
  - Make the same move twice, the opponent might not do the same thing
  - Rewards also stochastic (opponent does or doesn't take your piece)
- Temporal credit assignment problem
  - Did we get the reward because of this move? Or because we made good tactical decisions somewhere in the past?
- Exploration-exploitation trade-off
  - If we found a good opening, should we use this?
  - Or should we try other things, hoping for something better?



# Lecture 1 Overview

- What is machine learning ?
- Categories of machine learning
- **Notation**
- Approaching a machine learning application
- ML Terminology

# Supervised Learning Workflow -- Overview



# Supervised Learning Notation

- Training set:  $\mathcal{D} = \{\langle x^{[i]}, y^{[i]} \rangle, i = 1, \dots, n\}$ ,
- Unknown function:  $f(\mathbf{x}) = y$
- Hypothesis:  $h(\mathbf{x}) = \hat{y}$



**Classification**

$$h : \mathbb{R}^m \rightarrow \text{---}$$

**Regression**

$$h : \mathbb{R}^m \rightarrow \text{---}$$

# Data Representation

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix}$$

Feature vector

$$\mathbf{X} = \begin{bmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \vdots \\ \mathbf{x}_n^T \end{bmatrix}$$

Design matrix

$$\mathbf{X} = \begin{bmatrix} x_1^{[1]} & x_2^{[1]} & \cdots & x_m^{[1]} \\ x_1^{[2]} & x_2^{[2]} & \cdots & x_m^{[2]} \\ \vdots & \vdots & \ddots & \vdots \\ x_1^{[n]} & x_2^{[n]} & \cdots & x_m^{[n]} \end{bmatrix}$$

Design matrix

# Data Representation

$m =$  —

$n =$  —

**Samples**  
(instances, observations)

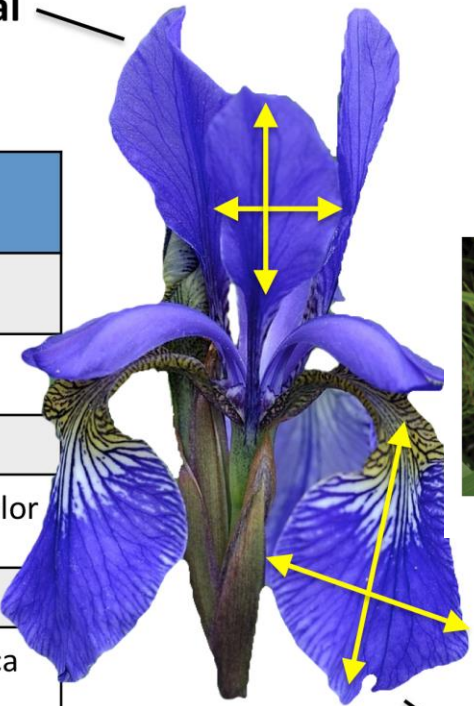
	Sepal length	Sepal width	Petal length	Petal width	Class label
1	5.1	3.5	1.4	0.2	Setosa
2	4.9	3.0	1.4	0.2	Setosa
...					
50	6.4	3.5	4.5	1.2	Versicolor
...					
150	5.9	3.0	5.0	1.8	Virginica

**Features**  
(attributes, measurements, dimensions)


**Class labels**  
(targets)

**Petal**

**Sepal**



Iris flower data set



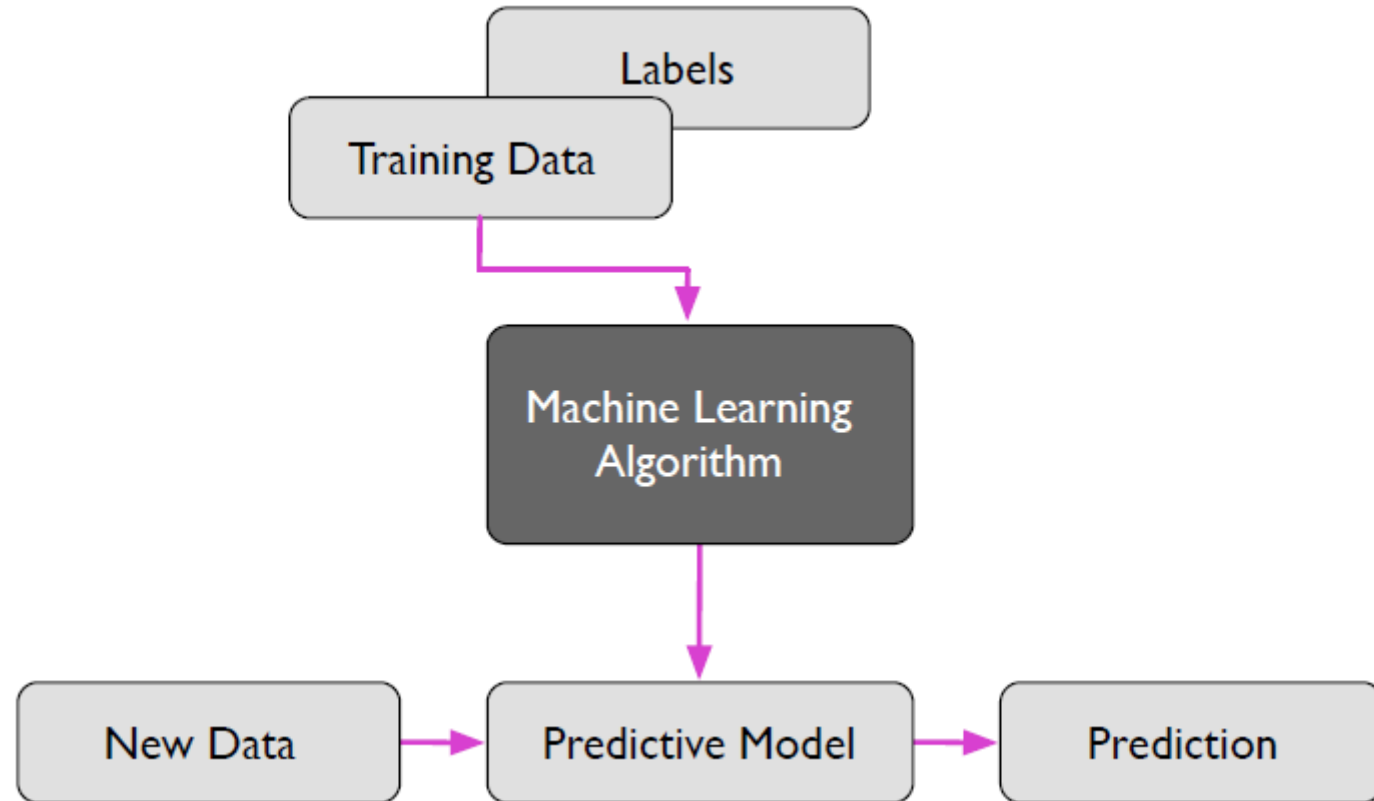
# ML Terminology

- **Training example:** A row in the table representing the dataset. Synonymous to an observation, training record, training instance, training sample.
- **Feature:** a column in the table representing the dataset. Synonymous to predictor, variable, input, attribute, covariate.
- **Targets:** What we want to predict. Synonymous to outcome, output, ground truth, response variable, dependent variable, (class) label (in classification).
- **Output / prediction:** use this to distinguish from targets; here, means output from the model.

# Lecture 1 Overview

- What is machine learning ?
- Categories of machine learning
- Notation
- Approaching a machine learning application

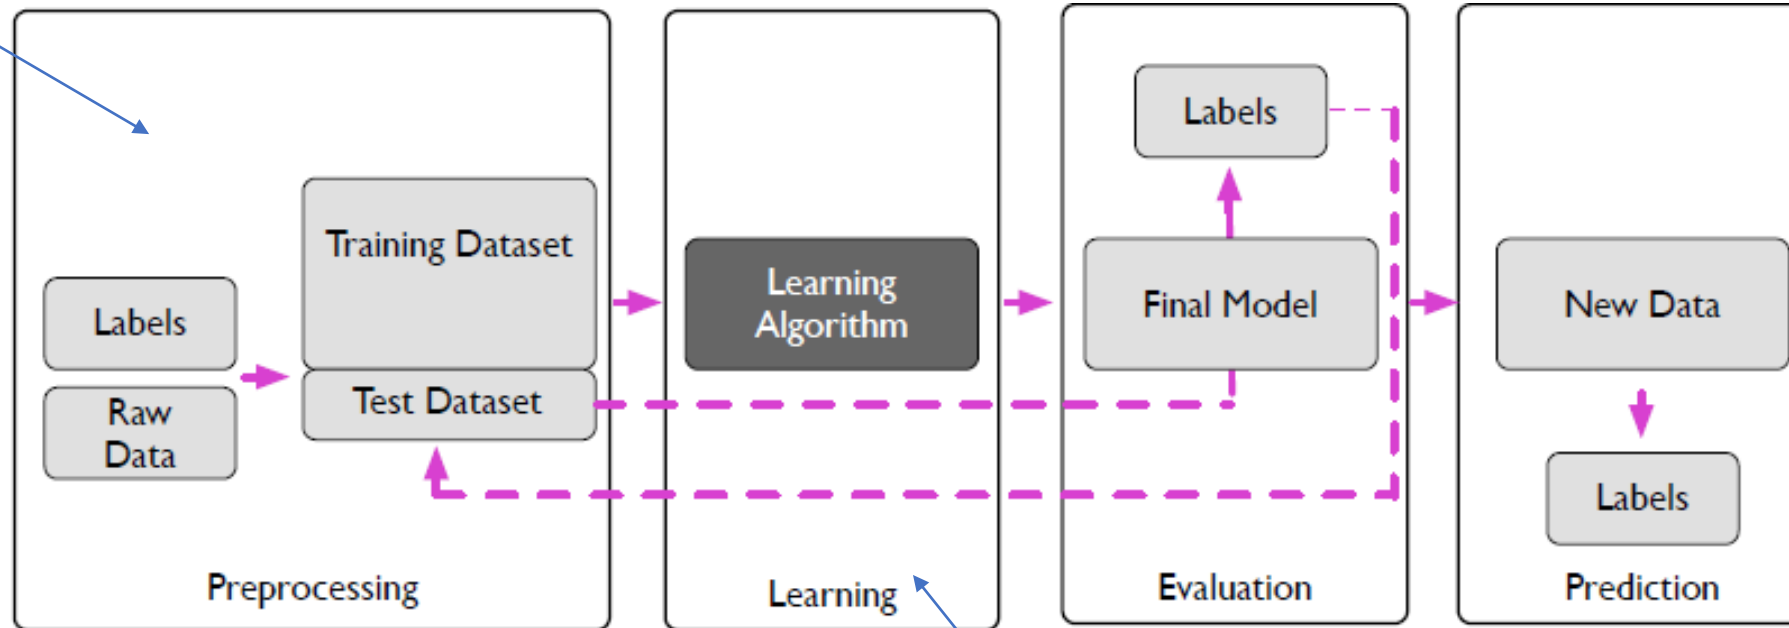
# Supervised Learning Workflow -- Overview





# The framework

Feature Extraction and Scaling  
Feature Selection  
Dimensionality Reduction  
Sampling



Model Selection  
Cross-Validation  
Performance Metrics  
Hyperparameter Optimization

# The framework

1. Define the problem to be solved.
2. Collect (labeled) data.
3. Choose an algorithm class.
4. Choose an optimization metric or measure for learning the model.
5. Choose a metric or measure for evaluating the model.

# Evaluation -- Misclassification Error

$$L(\hat{y}, y) = \begin{cases} 0 & \text{if } \hat{y} = y \\ 1 & \text{if } \hat{y} \neq y \end{cases}$$

$$ERR_{\mathcal{D}_{test}} = \frac{1}{n} \sum_{i=1}^n L(\hat{y}^i, y^i)$$

# ML Terminology

- **Loss function:** Often used synonymously with cost function; sometimes also called error function. In some contexts the loss for a single data point, whereas the cost function refers to the overall (average or summed) loss over the entire dataset. Sometimes also called empirical risk.